

Quiz 0

Assume that $y_n = \beta_0^* + x_n \beta_1^* + \varepsilon_n$ for scalar x_n and some fixed β_0^* and β_1^* . Assume that

- The residuals ε_n are IID with $\mathbb{E}[\varepsilon_n] = 0$ and $\text{Var}(\varepsilon_n) = \sigma^2$.
- The regressors x_n are IID with $\mathbb{E}[x_n] = \mu$ and $\text{Var}(x_n) = \nu^2$.
- The residuals are independent of the regressors.

You should be able to analyze quantities involving x_n , ε_n , and y_n in terms of the fixed quantities β_0^* , β_1^* , μ , σ , and ν .

- Be able to compute all the conditional variances and expectations you can think of
- Be able to apply the LLN and CLT, including recognizing when the CLT does not converge
- Be able to apply the LLN to random matrices and their inverses, including recognizing when the matrix is or is not invertible.

In addition, you should be able to

- Write out the $\mathbf{X}$, $\mathbf{Y}$, and $\boldsymbol{\varepsilon}$ matrices for simple linear regressions
- Write out the OLS estimator $\hat{\boldsymbol{\beta}}$ in terms of these quantities
- Be able to use the LLN and CLT to evaluate the limits of

$$\begin{aligned} & - \frac{1}{N} \mathbf{X}^\top \mathbf{X} \\ & - \frac{1}{N} \mathbf{X}^\top \mathbf{Y} \\ & - \frac{1}{\sqrt{N}} \mathbf{X}^\top \boldsymbol{\varepsilon} \end{aligned}$$

- Compute $\mathbb{E}[\hat{\boldsymbol{\beta}}|\mathbf{X}]$ and $\text{Cov}(\hat{\boldsymbol{\beta}}|\mathbf{X})$ in terms of $\text{Cov}(\boldsymbol{\varepsilon})$, and plug in for particular values of $\text{Cov}(\boldsymbol{\varepsilon})$.

In addition, you should be able to do all of the above for the special case when there is no x_n and there is only a constant.

Quiz 1

Suppose we observe data points of the form (y_n, s_n, a_n) , where a_n is a categorical variable and s_n is either a continuous variable or another categorical variable.

You should be able to

- Write the values of the $\mathbf{X}$ matrix for a one-hot encoding of a_n
- Write the values of the $\mathbf{X}$ matrix for interactions between s_n and a_n
- Understand the meaning of interactions (fitting a different slope for each level)

- Write the regression involving the one-hot encoding with and without a constant term, and understand how the interpretation of the coefficients change

Additionally, for all of these combinations, you should be able to do everything from Quiz 0.

Quiz 2

As a starting point, you should be able to do everything for Quiz 0 with a generic vector $\mathbf{x}$ rather than a set of one or two scalars.

You should be familiar with our four assumptions:

- Normality
- Homoskedasticity
- Heteroskedasticity
- Machine learning (IID)

I'll give you what these assumptions are, but you need to be able to understand and apply the definitions.

When we assume that $y_n = \beta^{*\top} \mathbf{x}_n + \varepsilon_n$, for each of our four assumptions, with enough hints you should be able to do everything from quiz 0, now in terms of β^* .

For testing, you should be able to analyze the t-statistic for a null of the form $\beta_k^* = 0$ for some component k .

- Know what the distribution of the t-statistic is under the null for normality (all you need is to understand its properties, you don't need to memorize anything)
- Know why you would use an F-statistic, and what its distribution is under the normal assumption (all you need is to understand its properties, you don't need to memorize anything)
- Know the distribution of the t-statistic for large N under non-normal assumptions
- Relatedly, understand how to compute the sandwich covariance and how it is different from the standard covariance
- Be able to compute the level of a given test by applying the above results under the null
- Be able to compute the power of a test against a particular alternative by plugging in an alternative β^*

The question will take the form of undertaking one or more of these tasks under a particular tractable choice of $\mathbf{x}$.

Quiz 3

This question will be about a one-dimensional example $y \sim z$ of a machine learning problem, focusing on squared loss.

You should be able to

- Write out the $\mathbf{X}$ matrix for a zeroth order spline basis of z
- Be able to plot the basis functions
- Understand intuitively what would change with higher-order splines (no need to know the exact formulas)
- Understand the relationship between the zeroth order spline, partitions of the domain of z , and one-hot encodings
- Know what the fitted value $\hat{y}$ is in each partition and be able to plot them (note: this is a generalization of quiz 1)
- Understand intuitively the costs and benefits of having more or fewer partitions

Additionally, you should be able to do all of Quiz 0 for ridge regression. You should also have an intuitive understanding of the costs and benefits of having a large or small ridge parameter.